

Xtend User Guide

December 6, 2011

Contents

1	Introduction	4
1.1	Getting Started	4
1.1.1	The Runtime Library	5
1.1.2	The Xtend Tutorial	5
2	Types	6
2.1	Common Super Type	6
2.2	Conformance and Conversion	6
3	Classes, Constructors, Fields and Methods	7
3.1	Package Declaration	7
3.2	Imports	7
3.3	Class Declaration	8
3.3.1	Inheritance	8
3.3.2	Examples	8
3.4	Constructors	8
3.5	Fields	9
3.6	Methods	10
3.6.1	Overriding Methods	10
3.6.2	Declared Exceptions	10
3.6.3	Inferred Return Types	11
3.6.4	Dispatch Methods	11
3.7	Annotations	14
3.8	Extension Methods	14
3.8.1	Local Methods	15
3.8.2	Library Extensions	15
3.8.3	Extension Imports	15
3.8.4	Extension Fields	15
4	Expressions	17
4.1	Literals	17
4.1.1	String Literals	17
4.1.2	Integer Literals	18
4.1.3	Boolean Literals	18
4.1.4	Null Literal	18
4.1.5	Type Literals	18
4.1.6	Function Types	18

4.2	Type Casts	19
4.3	Infix Operators and Operator Overloading	19
4.3.1	Short-Circuit Boolean Operators	20
4.3.2	Examples	21
4.3.3	Assignments	21
4.4	Variable Declarations	21
4.4.1	Typing	22
4.5	Feature Calls	22
4.5.1	Property Access	22
4.5.2	Implicit Variables <i>this</i> and <i>it</i>	23
4.5.3	Null-Safe Feature Call	23
4.6	Constructor Call	23
4.7	Closures	23
4.7.1	Typing	24
4.7.2	Function Mapping	24
4.7.3	Implicit Parameter <i>it</i>	25
4.7.4	Exceptions in Closures	25
4.7.5	Builder Syntax	25
4.8	If Expression	26
4.8.1	Typing	26
4.8.2	Examples	26
4.9	Switch Expression	27
4.9.1	Type guards	27
4.9.2	Typing	28
4.9.3	Examples	28
4.10	Blocks	28
4.10.1	Examples	28
4.11	For Loop	29
4.12	While Loop	29
4.12.1	Examples	29
4.13	Do-While Loop	29
4.13.1	Examples	30
4.14	Return Expression	30
4.15	Throwing Exceptions	30
4.16	Try, Catch, Finally	30
4.17	Template Expressions	31
4.17.1	Conditions in Templates	31
4.17.2	Loops in Templates	31
4.17.3	Typing	31
4.17.4	White Space Handling	32

1 Introduction

Xtend is a statically-typed programming language which is tightly integrated with and runs on the Java Virtual Machine. It has its roots in the Java programming language but improves on many concepts:

- *Advanced type inference* - You rarely need to write down type signatures
- *Full support for Java Generics* - Including all conformance and conversion rules
- *Closures* - concise syntax for anonymous function literals
- *Operator overloading* - make your libraries even more expressive
- *Powerful switch expressions* - type based switching with implicit casts
- *No statements* - Everything is an expression
- *Template expressions* - with intelligent white space handling
- *Extension methods* - Enhance closed types with new functionality
- *Property access syntax* - shorthands for getter and setter access
- *Multiple dispatch* a.k.a. polymorphic method invocation
- *Translates to Java* not bytecode - understand what is going on and use your code for platforms such as Android or GWT

It is not aiming at replacing Java all together. Therefore its library is a thin layer on top of the Java Development Kit (JDK) and interacts with Java exactly the same as it interacts with Xtend code. Java can call Xtend methods, too, in a completely transparent way. And of course, Xtend provides a modern Eclipse-based IDE closely integrated with the Java Development Tools (JDT).

1.1 Getting Started

describe how to get it

Let us start with a simple "Hello World" example. In Xtend, that reads as

```

class HelloWorld {
    def static main(String[] args) {
        println("Hello World")
    }
}

```

The Xtend code resembles Java a lot. You can already see how the syntactic noise is reduced: No semicolons, no return types etc.

An Xtend class has to reside in a Java project which allows to setup its classpath. If the project has the Xtext nature, the Xtend class will automatically be compiled to Java code in the *xtend-gen* source folder. The resulting Java code will look like this

```

import org.eclipse.xtext.xbase.lib.InputOutput;

@SuppressWarnings("all")
public class HelloWorld {
    public static String main(final String[] args) {
        String _println = InputOutput.<String>println("Hello World");
        return _println;
    }
}

```

1.1.1 The Runtime Library

The only surprising fact in the generated Java code may be the library class `InputOutput`. Many features of Xtend are not built into the language itself but provided via small libraries *org.eclipse.xtend.lib.jar*. This library and its dependencies have to reside on the classpath of each Xtend class. The libraries provide a lot of useful helper methods to create and manipulate collections, default operator bindings for primitives and other handy things the generated Java code relies on.

1.1.2 The Xtend Tutorial

The best way to get acquainted with the language is to materialize the *Xtend Tutorial* example project in your workspace. You will find it in the new project wizard dialog.

The project contains a couple of sample Xtend files, which show the different language concepts in action. Looking into the *xtend-gen* folder which holds the compiled Java code will help you understand the concepts better.

2 Types

Xtend fully supports Java's type system: The primitive types as **int** or **boolean** are available as well as all classes and interfaces that reside on the classpath.

Java Generics are fully integrated, such that you can define type parameters and type arguments in just the same way as in Java.

As Xtend classes compile to Java classes, you can integrate both Java classes in Xtend and Xtend classes in Java. Note that Xtend does not have an own syntax for interfaces, as the Java syntax is already very concise.

2.1 Common Super Type

Because of type inference Xbase sometimes needs to compute the most common super type of a given set of types.

For a set $[T1, T2, \dots, Tn]$ of types the common super type is computed by using the linear type inheritance sequence of $T1$ and is iterated until one type conforms to each $T2, \dots, Tn$. The linear type inheritance sequence of $T1$ is computed by ordering all types which are part of the type hierarchy of $T1$ by their specificity. A type $T1$ is considered more specific than $T2$ if $T1$ is a subtype of $T2$. Any types with equal specificity will be sorted by the maximal distance to the originating subtype. *CharSequence* has distance 2 to *StringBuilder* because the supertype *AbstractStringBuilder* implements the interface, too. Even if *StringBuilder* implements *CharSequence* directly, the interface gets distance 2 in the ordering because it is not the most general class in the type hierarchy that implements the interface. If the distances for two classes are the same in the hierarchy, their qualified name is used as the compare-key to ensure deterministic results.

2.2 Conformance and Conversion

Conformance is used in order to find out whether some expression can be used in a certain situation. For instance when assigning a value to a variable, the type of the right hand expression needs to conform to the type of the variable.

As Xbase implements the unchanegd type system of Java it also fully supports the conformance rules defined in The Java Language Specification.

3 Classes, Constructors, Fields and Methods

At a first glance an Xtend file pretty much looks like a Java file. It starts with a package declaration followed by an import section and a class definition. The class in fact is directly translated to a Java class in the corresponding Java package. As in Java, a class can have constructors, fields and methods.

Here is an example:

```
package com.acme

import java.util.List

class MyClass {
    String name

    new(String name) {
        this.name = name
    }

    def String first(List<String> elements) {
        elements.get(0)
    }
}
```

3.1 Package Declaration

Package declarations are mostly like in Java. There are two the small differences:

- An identifier can be escaped with a ^ character in case it conflicts with a keyword.
- There is no terminating semicolon.

```
package org.eclipse.xtext
```

3.2 Imports

The ordinary imports of type names are equivalent to the imports known from Java. Again one can escape any names conflicting with keywords using a ^. In contrast to

Java, the import statement is never terminated with a semicolon. Xtend also features static imports but allows only a wildcard `*` at the end, i.e. you cannot import single members using a static import. For non-static imports the use of wildcards is deprecated for the benefit of better tooling.

As in Java all classes from the `java.lang` package are implicitly imported.

```
import java.math.BigDecimal
import static java.util.Collections.*
```

Static methods of helper classes can also be imported as *extensions*. See the section on extension methods (§3.8) for details.

3.3 Class Declaration

The class declaration reuses a lot of Java's syntax but still is a bit different in some aspects: Java's default "package private" visibility does not exist in Xtend. As an Xtend class is compiled to a top-level Java class and Java does not allow **private** or **protected** top-level classes any Xtend class is **public**. It is possible to write this explicitly.

To be implemented: The **abstract** as well as the **final** modifiers are directly translated to Java and have the exact same meaning.

3.3.1 Inheritance

Inheritance is directly reused from Java. Single inheritance of Java classes as well as implementing multiple Java interfaces is supported. Because Xtend classes are compiled to Java, Xtend classes can extend other Xtend classes, and even Java classes can inherit from Xtend classes.

3.3.2 Examples

The most simple class :

```
class MyClass {
}
```

A more advanced class declaration in Xtend :

```
class ArrayList<E> extends AbstractList<E>
    implements List<E>, RandomAccess,
        Cloneable, java.io.Serializable {
    ...
}
```

3.4 Constructors

An Xtend class can define one or more constructors. Unlike Java, constructors are always named *new*. Constructors can also delegate to other constructors using **this**(args...) in their first line.


```
class Foo {
    new(String s) {
        ...
    }

    new() {
        this("default")
    }
    ...
}
```

The same rules with regard to inheritance apply as in Java, i.e. if the super class does not define a no-argument constructor, you have to explicitly call one using **super**(args...) as the first expression in the body of the constructor:

```
class Foo extends Bar {
    new(String s) {
        super(s.length)
    }
    ...
}
```

The default visibility of constructors is **public** but you can also specify **protected** or **private**.

3.5 Fields

An Xtend class can define fields, too. The syntax is the same as in Java. Fields can be declared and initialized in the same line. Fields marked as **static** will be compiled to static Java fields. Any other modifier from Java on fields is not yet supported.

```
class Foo {
    String name
    int count = 1
    Foo foo = new Foo()
    static boolean debug = false
}
```

The default visibility is **private**. You can also declare it explicitly as being **public**, **protected**, or **private**. Alternatively, you can write accessor methods.

Fields in Xtend are commonly used together with an annotation for a dependency injection container. Example:

```
@Inject MyService myService
```

This will translate to the following Java field:

```
@Inject
private MyService myService;
```

A speciality of Xtend are *extension methods* which are covered in their own section (§3.8).

3.6 Methods

Xtend methods are declared within a class and are translated to a corresponding Java method with exactly the same signature. The only exceptions are dispatch methods, which are explained here (§3.6.4).

```
def boolean equalsIgnoreCase(String s1, String s2) {  
    s1.toLowerCase() == s2.toLowerCase();  
}
```

The default visibility of a plain method is **public**. You can explicitly declare it as being **public**, **protected**, or **private**.

Xtend supports the **static** modifier for methods:

```
def static createInstance() {  
    new Foo()  
}
```

3.6.1 Overriding Methods

Methods can override other methods from the super class or implemented interface methods using the keyword **override**. If a method overrides a method from a super type, the override keyword is mandatory and replaces the keyword **def**. As in Java **final** methods cannot be overridden by subclasses.

Example:

```
override boolean equalsIgnoreCase(String s1,String s2) :  
    s1.toLowerCase() == s2.toLowerCase();
```

3.6.2 Declared Exceptions

Xtend does not force you to catch checked exceptions. Instead, they are rethrown in a way the compiler does not complain about a missing throws clause, using the sneaky-throw technique introduced by Lombok. Nevertheless, you can still declare the exceptions thrown in a method's body using the same **throws** clause as in Java.

```
/*  
 * throws an Exception  
 */  
def void throwException() throws Exception {  
    throw new Exception()  
}  
  
/*  
 * throws an Exception without declaring it
```

```

*/
def void sneakyThrowException() {
    throw new Exception()
}

```

3.6.3 Inferred Return Types

If the return type of a method can be inferred from its body it does not have to be declared. That is the method

```

def boolean equalsIgnoreCase(String s1,String s2) :
    s1.toLowerCase() == s2.toLowerCase();

```

could be declared like this:

```

def equalsIgnoreCase(String s1,String s2) :
    s1.toLowerCase() == s2.toLowerCase();

```

This does not work for abstract method declarations as well as if the return type of a method depends on a recursive call of the same method. The compiler tells the user when it needs to be specified.

3.6.4 Dispatch Methods

Generally method binding works just like method binding in Java. That is method calls are bound based on the static types of arguments. Sometimes this is not what you want. Especially in the context of extension methods (§3.8) you would like to have polymorphic behavior.

Dispatch methods make a set of overloaded methods polymorphic. That is the runtime types of all given arguments are used to decide which of the overloaded methods is being invoked. This essentially removes the need for the quite invasive visitor pattern.

A dispatch method is marked using the keyword **dispatch**.

```

def dispatch foo(Number x) { "it's a number" }
def dispatch foo(Integer x) { "it's an int" }

```

For a set of visible dispatch methods in the current type hierarchy, the compiler infers a common signature using the common super type of all declared arguments and generates a Java method made up of **if instanceof else** cascades. It dispatches to the different available methods. The actually declared methods are all compiled to Java methods that are prefixed with an underscore.

For the two dispatch methods in the example above the following Java code would be generated:

```

public String foo(Number x) {
    if (x instanceof Integer) {
        return _foo((Integer)x);
    } else if (x instanceof Number) {
        return _foo((Number)x);
    } else {

```

```

        throw new IllegalArgumentException(
            "Couldn't handle argument x:" + x);
    }
}

protected String _foo(Integer x) {
    return "It's an int";
}

protected String _foo(Number x) {
    return "It's a number";
}

```

Note that the **instanceof** cascade is ordered by how specific a type is. More specific types come first.

The default visibility of dispatch methods is **protected**. If all dispatch methods explicitly declare the same visibility, this will be the visibility of the inferred dispatcher, too. Otherwise it is **public**

In case there is no single most general signature, one is computed and the different overloaded methods are matched in the order they are declared within the class file. Example:

```

def dispatch foo(Number x, Integer y) { "it's some number and an int" }
def dispatch foo(Integer x, Number x) { "it's an int and a number" }

```

generates the following Java code :

```

public String foo(Number x, Number y) {
    if ((x instanceof Number)
        && (y instanceof Integer)) {
        return _foo((Number)x, (Integer)y);
    } else if ((x instanceof Integer)
        && (y instanceof Number)) {
        return _foo((Integer)x, (Number)y);
    } else {
        throw new IllegalArgumentException(
            "Couldn't handle argument x:" + x + ", argument y:" + y);
    }
}

```

As you can see a **null** reference is never a match. If you want to fetch **null** you can declare a parameter using the type `java.lang.Void`.

```

def dispatch foo(Number x) { "it's some number" }
def dispatch foo(Integer x) { "it's an int" }
def dispatch foo(Void x) { throw new NullPointerException("x") }

```

Which compiles to the following Java code:

```

public String foo(Number x) {
    if (x instanceof Integer) {

```

```

        return _foo((Integer)x);
    } else if (x instanceof Number){
        return _foo((Number)x);
    } else if (x == null) {
        return _foo((Void)null);
    } else {
        throw new IllegalArgumentException(
            "Couldn't handle argument x:" + x);
    }
}

```

Dispatch Methods and Inheritance

Any visible Java methods from super types conforming to the compiled form of a dispatch method are also included in the dispatch. Conforming means they have the right number of arguments and have the same name (starting with an underscore).

For example, consider the following Java class :

```

public abstract class AbstractLabelProvider {
    protected String _label(Object o) {
        // some generic implementation
    }
}

```

and the following Xtend class which extends the Java class :

```

class MyLabelProvider extends AbstractLabelProvider {
    def dispatch label(Entity it) {
        name
    }

    def dispatch label(Method it) {
        name+"(" + params.toString(",") + "):" + type
    }

    def dispatch label(Field it) {
        name + type
    }
}

```

The resulting dispatch method in the generated Java class 'MyLabelProvider' would then look like this:

```

public String label(Object o) {
    if (o instanceof Field) {
        return _label((Field)o);
    } else if (o instanceof Method){
        return _foo((Method)o);
    } else if (o instanceof Entity){
        return _foo((Entity)o);
    }
}

```

```

    } else if (o instanceof Object){
        return _foo((Object)o);
    } else {
        throw new IllegalArgumentException(
            "Couldn't handle argument o:" + o);
    }
}

```

Static Dispatch Methods

Even static dispatch methods are allowed. The same rules apply, but you cannot mix static and non-static dispatch methods.

3.7 Annotations

Xtend supports Java annotations. The syntax is exactly like defined in the Java Language Specification. Annotations are available on classes, fields, methods and parameters.

Example:

```

@TypeAnnotation(typeof(String))
class MyClass {
    @FieldAnnotation(children = {@MyAnno(true), @MyAnnot(false)})
    String myField

    @MethodAnnotation(children = {@MyAnno(true), @MyAnnot})
    def String myMethod(@ParameterAnnotation String param) {
        //...
    }
}

```

3.8 Extension Methods

Extensions methods are a technique to add methods to existing classes without modifying their code. This feature is actually where Xtend has its name from. They are based on a simple syntactic trick: Instead of passing the first argument of an extension method inside the parentheses of a call, the method is called on the argument parameter as if it was one of its members.

```

"foo".toFirstUpper() // calls toFirstUpper("foo")

```

Method calls in extension syntax often result in much better readable code, as function calls are rather concatenated than nested. They also allow to add methods in a certain context only.

To be callable in the extension syntax, a method has to be on the extension scope of the callers context. Xtend supports a variety of ways to put methods on this scope.

3.8.1 Local Methods

All methods of the current Xtend class are automatically available in extension syntax. For example

```
class Foo {  
    def foo(Bar bar) {  
        // do something with bar  
    }  
  
    def extensionCall(Bar bar) {  
        bar.foo() // calls this.foo(bar)  
    }  
}
```

3.8.2 Library Extensions

The static methods of the classes in the Xtend runtime library (§1.1.1) are automatically available as extensions, e.g.

```
newArrayList() // CollectionLiterals.newArrayList()  
"foo".toFirstUpper // StringExtensions.toFirstUpper(String)
```

3.8.3 Extension Imports

In Java, you would usually write a helper class with static helper functions to decorate an existing class with additional behavior. In order to integrate such static helper classes, Xtend allows to put the keyword **extension** after the **static** keyword of a static import (§3.2) thus putting making all imported static functions available as extension methods.

That is the following import declaration

```
import static extension java.util.Collections.*
```

allows to use its methods for example like this :

```
new Foo().singletonList() // calls Collections.singletonList(new Foo)
```

Although this is supported it is generally much nicer to use extension fields (§3.8.4), because they allow to change the bound implementation.

3.8.4 Extension Fields

You can make the instance methods provided by the field available as extension methods, by adding the keyword **extension** to the field declaration.

Imagine you want to add a method *fullName* to a closed type *Entity*. With extension methods, you could declare the following class

```
class PersonExtensions {  
    def getFullName(Person p) {  
        p.forename + " " + p.name  
    }  
}
```

```
}  
}
```

And if you have an instance of this class injected as extension like this:

```
class PersonPrinter {  
    @Inject extension PersonExtensions  
    ...  
}
```

The method is being put on the extension scope of the class `Person`. This is why you can skip the field's name. You can now write the following

```
def print(Person myPerson) {  
    myPerson.getFullName()  
}
```

which is translated to the Java method

```
@Override  
public String print(Person myPerson) {  
    return _personExtensions.getFullName(myPerson);  
}
```

where `_personExtensions` is the default name of the field. Of course the property shorthand (see section on property access (§4.5.1)) is still available.

```
myPerson.fullName
```

The nice thing with using dependency injection in combination with the extension modifier as opposed to extension imports (§3.8.3) is, that in case there is a bug in the extension or it is implemented inefficiently or you just need a different strategy, you can simply exchange the component with another implementation. You do this without modifying the library nor the client code. You will only have to change the binding in your DI configuration. Also this gives you a general hook for any AOP-like thing you would want to do, or allows you to write against an SPI, where the concrete implementation can be provided by a third party.

4 Expressions

There are no statements in Xtend. Instead, everything is an expression and has a return value. That allows to use every Xtend expression on the right hand side of an assignment. For example, as a

```
try  
catch
```

is an expression the following code is legal in Xtend:

```
val data = try {  
    fileContentsToString('data.txt')  
} catch (IOException e) {  
    'dummy data'  
}
```

If `fileContentsToString()` throws an `IOException`, it is caught and a default value is returned and thus assigned to the variable `data`.

In Xtend, expressions appear as initializers of fields (§3.5) or as the bodies of constructors or methods. A method body in Xtend can either be a single block expression (§4.10) or a template expression (§4.17).

4.1 Literals

A literal denotes a fixed unchangeable value. Literals for strings, integers, boolean values, **null** and Java types are supported.

4.1.1 String Literals

A string literal is a valid expression and returns an instance of `java.lang.String` of the given value. String literals are enclosed by a pair of single quotes or double quotes allowing to use the respective other unquoted inside the string. Special characters can be quoted with a backslash or defined using Java's unicode notation. Strings can span multiple lines.

- 'Hello World !'
- "Hello World !"
- 'Hello "World" !'
- "Hello \"World\" !"'

- `'\u00a1Hola el mundo!'`
- `"Hello
World !"`

4.1.2 Integer Literals

An integer literal creates an **int**. There is no signed integer. If you put a minus operator in front of an integer literal it is taken as a `UnaryOperator` with one argument (the positive integer literal).

- `42`
- `234254`

4.1.3 Boolean Literals

There are two boolean literals, **true** and **false** which correspond to their Java counterpart of type **boolean**.

- **true**
- **false**

4.1.4 Null Literal

The null pointer literal is, like in Java, **null**. It is a member of any reference type and the only member of the type `java.lang.Void`.

- **null**

4.1.5 Type Literals

Type literals are specified using the keyword **typeof** :

- **typeof**(`java.lang.String`) which yields `java.lang.String.class`

4.1.6 Function Types

Xbase introduces closures (§4.7), and therefore an additional function type signature. On the JVM-Level a closure (or more generally any function object) is just an instance of one of the types in `org.eclipse.xtext.xbase.lib.Function*`, depending on the number of arguments. However, as closures are a very important language feature, a special sugared syntax for function types has been introduced. So instead of writing `Function1<String,Boolean>` one can write `(String)=>Boolean`.

Primitives cannot be used in function types.

4.2 Type Casts

Type cast behave like casts in Java, but have a slightly better readable syntax. Type casts bind stronger than any other operator but weaker than feature calls.

The conformance rules for casts are defined in the Java Language Specification. Here is an example:

- `my.foo as MyType`

4.3 Infix Operators and Operator Overloading

There are a couple of common predefined infix operators. In contrast to Java, the operators are not limited to operations on certain types. Instead an operator-to-method mapping allows users to redefine the operators for any type just by implementing the corresponding method signature. As an example, the Xtend runtime library (§1.1.1) contains a class `BigDecimalExtensions` that defines operators for `BigDecimal`s which allows the following code:

```
val x=new BigDecimal('2.71')
val y=new BigDecimal('3.14')
x+y // calls BigDecimalExtension.operator_plus(x,y)
```

The following defines the operators and the corresponding Java method signatures / expressions.

<code>e1 += e2</code>	<code>e1.operator_add(e2)</code>
<code>e1 e2</code>	<code>e1.operator_or(e2)</code>
<code>e1 && e2</code>	<code>e1.operator_and(e2)</code>
<code>e1 == e2</code>	<code>e1.operator_equals(e2)</code>
<code>e1 != e2</code>	<code>e1.operator_notEquals(e2)</code>
<code>e1 < e2</code>	<code>e1.operator_lessThan(e2)</code>
<code>e1 > e2</code>	<code>e1.operator_greaterThan(e2)</code>
<code>e1 <= e2</code>	<code>e1.operator_lessEqualsThan(e2)</code>
<code>e1 >= e2</code>	<code>e1.operator_greaterEqualsThan(e2)</code>
<code>e1 -> e2</code>	<code>e1.operator_mappedTo(e2)</code>
<code>e1 .. e2</code>	<code>e1.operator_upTo(e2)</code>
<code>e1 + e2</code>	<code>e1.operator_plus(e2)</code>
<code>e1 - e2</code>	<code>e1.operator_minus(e2)</code>
<code>e1 * e2</code>	<code>e1.operator_multiply(e2)</code>
<code>e1 / e2</code>	<code>e1.operator_divide(e2)</code>
<code>e1 % e2</code>	<code>e1.operator_modulo(e2)</code>
<code>e1 ** e2</code>	<code>e1.operator_power(e2)</code>
<code>! e1</code>	<code>e1.operator_not()</code>
<code>- e1</code>	<code>e1.operator_minus()</code>

The table above also defines the operator precedence in ascending order. The blank lines separate precedence levels. The assignment operator `+=` is right-to-left associative in the same way as the plain assignment operator `=` is. That is `a = b = c` is executed as `a = (b = c)`, all other operators are left-to-right associative. Parenthesis can be used to adjust the default precedence and associativity.

4.3.1 Short-Circuit Boolean Operators

If the operators `||` and `&&` are used in a context where the left hand operand is of type boolean, the operation is evaluated in short circuit mode, which means that the right hand operand might not be evaluated at all in the following cases:

1. in the case of `||` the operand on the right hand side is not evaluated if the left operand evaluates to **true**.
2. in the case of `&&` the operand on the right hand side is not evaluated if the left operand evaluates to **false**.

4.3.2 Examples

- `my.foo = 23`
- `myList += 23`
- `x > 23 && y < 23`
- `x && y || z`
- `1 + 3 * 5 * (- 23)`
- `!(x)`
- `my.foo = 23`
- `my.foo = 23`

4.3.3 Assignments

Local variables (§4.4) can be reassigned using the `=` operator. Also properties can be set using this operator: Given the expression

```
myObj.myProperty = "foo"
```

The compiler first looks up whether there is an accessible Java Field called `myProperty` on the type of `myObj`. If there is one it translates to the following Java expression :

```
myObj.myProperty = "foo";
```

Remember in Xtend everything is an expression and has to return something. In the case of simple assignments the return value is the value returned from the corresponding Java expression, which is the assigned value.

If there is no accessible field on the left operand's type, a method called `setMyProperty` (`OneArg`) (JavaBeans setter method) is looked up. It has to take one argument of the type (or a super type) of the right hand operand. The return value will be whatever the setter method returns (which usually is null). As a result the compiler translates to :

```
myObj.setMyProperty("foo")
```

4.4 Variable Declarations

Variable declarations are only allowed within blocks (§4.10). They are visible in any subsequent expressions in the block.

A variable declaration starting with the keyword **val** denotes a so called value, which is essentially a final (i.e. unsettable) variable. In rare cases, one needs to update the value of a reference. In such situations the variable needs to be declared with the keyword **var**, which stands for 'variable'.

A typical example for using **var** is a counter in a loop.

```

{
    val max = 100
    var i = 0
    while (i < max) {
        println("Hi there!")
        i = i + 1
    }
}

```

Although overriding or shadowing variables from outer scopes is allowed, it is usually only used to overload the implicit variable `it` (§4.5.2), in order to subsequently access an object's features in an unqualified manner.

Variables declared outside a closure using the `var` keyword are not accessible from within a closure.

4.4.1 Typing

The return type of a variable declaration expression is always `void`. The type of the variable itself can either be explicitly declared or be inferred from the right hand side expression. Here is an example for an explicitly declared type:

```
var List<String> msg = new ArrayList<String>();
```

In such cases, the right hand expression's type must conform (§2.2) to the type on the left hand side.

Alternatively the type can be left out and will be inferred from the initialization expression:

```
var msg = new ArrayList<String>(); // -> type ArrayList<String>
```

4.5 Feature Calls

A feature call is used to invoke members of objects, such as fields and methods, but also can refer to local variables and parameters, which are made available for the current expression's scope.

4.5.1 Property Access

Feature calls are directly translated to their Java equivalent with the exception, that for calls to properties an equivalent rule as described in subsection 4.3.3 applies. That is, for the following expression

```
myObj.myProperty
```

the compiler first looks for an accessible field in the type of `myObj`. If no such field exists it looks for a method called `myProperty()` before it looks for the getter methods `getMyProperty()`. If none of these members can be found the expression is unbound and a compilation error is thrown.

4.5.2 Implicit Variables `this` and `it`

If the current scope contains a variable named `this` or `it`, the compiler will make all its members available to the scope. That is one of

```
it.myProperty  
this.myProperty
```

is a valid expression

```
myProperty
```

is valid as well and is equivalent, as long as there is no local variable 'myProperty' on the scope, which would have higher precedence.

As `this` is bound to the surrounding object in Java, `it` can be used in finer-grained constructs such as function parameters. That is why `it.myProperty` has higher precedence than `this.myProperty`. `it` is also the default parameter name in closures (§4.7.3).

4.5.3 Null-Safe Feature Call

Checking for null references can make code very unreadable. In many situations it is ok for an expression to return null if a receiver was null. Xtend supports the safe navigation operator `?.` to make such code more readable.

Instead of writing

```
if ( myRef != null ) myRef.doStuff()
```

one can write

```
myRef?.doStuff()
```

4.6 Constructor Call

Construction of objects is done by invoking Java constructors. The syntax is exactly as in Java, e.g.

- `new String()`
- `new java.util.ArrayList<java.math.BigDecimal>()`

4.7 Closures

A closure is a literal that defines an anonymous function. A closure also captures the current scope, so that any final variables and parameters visible at construction time can be referred to in the closure's expression.

The surrounding square brackets are optional if the closure is the single argument of a method invocation. That is you can write

```
myList.find(e|e.name==null)
```

instead of

```
myList.find([e|e.name==null])
```

But in all other cases the square brackets are mandatory:

```
val func = [String s| s.length>3]
```

4.7.1 Typing

Closures are expressions which produce *Function* objects. The type is a function type (§4.1.6), consisting of the types of the parameters as well as the return type. The return type is never specified explicitly but is always inferred from the expression. The parameter types can be inferred if the closure is used in a context where this is possible.

For instance, given the following Java method signature:

```
public T <T>getFirst(List<T> list, Function0<T,Boolean> predicate)
```

the type of the parameter can be inferred. Which allows users to write:

```
arrayList( "Foo", "Bar" ).findFirst( e | e == "Bar" )
```

instead of

```
arrayList( "Foo", "Bar" ).findFirst( String e | e == "Bar" )
```

Here are some more examples:

- [| "foo"] //closure without parameters
- [String s | s.toUpperCase()] //explicit argument type
- [a,b,a | a+b+c] //inferred argument types

4.7.2 Function Mapping

An Xtend closure is a Java object of one of the *Function* interfaces shipped with the runtime library of Xtend. There is an interface for each number of parameters (current maximum is six parameters). The names of the interfaces are

- *Function0<ReturnType>* for zero parameters,
- *Function1<Param1Type, ReturnType>* for one parameters,
- *Function2<Param1Type, Param2Type, ReturnType>* for two parameters,
- ...
- *Function6<Param1Type, Param2Type, Param3Type, Param4Type, Param5Type, Param6Type, ReturnType>* for six parameters,

or

- *Procedure0* for zero parameters,
- *Procedure1<Param1Type>* for one parameters,
- *Procedure2<Param1Type, Param2Type>* for two parameters,
- ...
- *Procedure6<Param1Type, Param2Type, Param3Type, Param4Type, Param5Type, Param6Type>* for six parameters,

if the return type is **void**.

In order to allow seamless integration with existing Java libraries such as the JDK or Google Guava (formerly known as Google Collect) closures are automatically coerced to expected types if those types declare only one method (methods from `java.lang.Object` do not count).

As a result given the method `java.util.Collections.sort(List<T>, Comparator<? super T>)` is available as an extension method, it can be invoked like this

```
newArrayList( 'aaa', 'bb', 'c' ).sort(
    e1, e2 | if ( e1.length > e2.length ) {
        -1
    } else if ( e1.length < e2.length ) {
        1
    } else {
        0
    })
```

4.7.3 Implicit Parameter *it*

If a closure has a single parameter whose type can be inferred, the declaration of the parameter can be omitted. Use *it* to refer to the parameter inside the closure's body.

```
val (String s)=>String function = [toUpperCase]
// equivalent to [it | it.toUpperCase]
```

4.7.4 Exceptions in Closures

Checked exceptions that are thrown in the body of a closure are rethrown using the sneaky-throw technique (§3.6.2), i.e. you do not have to declare them explicitly.

4.7.5 Builder Syntax

If the last argument of a function call is a closure, it can be appended after the parenthesized parameter list. In combination with the implicit **it** parameter, skipping empty parentheses, and extension methods, this yields a very concise syntax.

```
val fruit = newArrayList('apple', 'pear', 'lemon')
fruit.map[toUpperCase]
// same as map(fruit, [it | it.toUpperCase])
```

This feature is especially useful when you are building object trees. A common pattern is to provide a set of extension functions taking two parameters: the parent object and a closure to initialize the new child. Here is an example for creating a simple tree of Nodes:

```
def createNode(Node parent, (Node)=>Void initializer) {
    val child=new Node()
    initializer.apply(child)
    child
}

newNode(null) [
    label="root"
    children += createNode [
        label="child0"
    ]
    children += createNode [
        label="child1"
    ]
]
```

4.8 If Expression

An if expression is used to choose two different values based on a predicate. While it has the syntax of Java's if statement it behaves like Java's ternary operator (`predicate ? thenPart : elsePart`), i.e. it is an expression that returns a value. Consequently, you can use if expressions deeply nested within expressions.

An expression `if (p)e1 else e2` results in either the value `e1` or `e2` depending on whether the predicate `p` evaluates to `true` or `false`. The else part is optional which is a shorthand for `else null`. That means

```
if (foo) x
```

is the a short hand for

```
if (foo) x else null
```

4.8.1 Typing

The type of an if expression is calculated by the return types `T1` and `T2` of the two expression `e1` and `e2`. It uses the rules defined in section 2.1.

4.8.2 Examples

- `if (isFoo) this else that`
- `if (isFoo){ this } else if (thatFoo){ that } else { other }`
- `if (isFoo) this`

4.9 Switch Expression

The switch expression is a bit different from Java's. First, there is no fall through which means only one case is evaluated at most. Second, the use of switch is not limited to certain values but can be used for any object reference instead.

For a switch expression

```
switch e {  
  case e1 : er1  
  case e2 : er2  
  ...  
  case en : ern  
  default : er  
}
```

the main expression `e` is evaluated first and then each case sequentially. If the switch expression contains a variable declaration using the syntax known from section 4.11, the value is bound to the given name. Expressions of type `java.lang.Boolean` or `boolean` are not allowed in a switch expression.

The guard of each case clause is evaluated until the switch value equals the result of the case's guard expression or if the case's guard expression evaluates to `true`. Then the right hand expression of the case evaluated and the result is returned.

If none of the guards matches the default expression is evaluated and returned. If no default expression is specified the expression evaluates to `null`.

Example:

```
switch myString {  
  case myString.length>5 : 'a long string.'  
  case 'foo' : 'It's a foo.'  
  default : 'It's a short non-foo string.'  
}
```

4.9.1 Type guards

In addition to the case guards one can add a so called *Type Guard* which is syntactically just a type reference (§2) preceding the optional case keyword. The compiler will use that type for the switch expression in subsequent expressions. Example:

```
var Object x = ...;  
switch x {  
  String case x.length()>0 : x.length()  
  List<?> : x.size()  
  default : -1  
}
```

Only if the switch value passes a type guard, i.e. an `instanceof` operation returns `true`, the case's guard expression is executed using the same semantics explained in previously.

If the switch expression contains an explicit declaration of a local variable or the expression references a local variable, the type guard acts like a cast, that is all references to the switch value will be of the type specified in the type guard.

4.9.2 Typing

The return type of a switch expression is computed using the rules defined in section 2.1. The set of types from which the common super type is computed corresponds to the types of each case's result expression. In case a switch expression's type is computed using the expected type from the context, it is sufficient to return the expected type if all case branches types conform to the expected type.

4.9.3 Examples

- ```
switch foo {
 Entity : foo.superType.name
 Datatype : foo.name
 default : throw new IllegalStateException
}
```
- ```
switch x : foo.bar.complicated('hello',42) {  
    case "hello42" : ...  
    case x.length<2 : ...  
    default : ....  
}
```

4.10 Blocks

The block expression allows to have imperative code sequences. It consists of a sequence of expressions, and returns the value of the last expression. The return type of a block is also the type of the last expression. Empty blocks return null. Variable declarations (§4.4) are only allowed within blocks and cannot be used as a block's last expression.

A block expression is surrounded by curly braces and contains at least one expression. It can optionally be terminated by a semicolon.

4.10.1 Examples

```
{  
    doSideEffect("foo")  
    result  
}
```

```
{  
    var x = greeting();  
    if (x.equals("Hello ")) {  
        x+"World!";  
    } else {  
        x;  
    }  
}
```

```
}  
}
```

4.11 For Loop

The for loop **for** (T1 variable : iterableOfT1)expression is used to execute a certain expression for each element of an array of an instance of `java.lang.Iterable`. The local variable is final, hence cannot be updated.

The return type of a for loop is **void**. The type of the local variable can be left out. In that case it is inferred from the type of the array or `java.lang.Iterable` returned by the iterable expression.

- **for** (String s : myStrings) {
 doSideEffect(s);
}

- **for** (s : myStrings)
 doSideEffect(s)

4.12 While Loop

A while loop **while** (predicate)expression is used to execute a certain expression unless the predicate is evaluated to **false**. The return type of a while loop is **void**.

4.12.1 Examples

- **while** (**true**) {
 doSideEffect("foo");
}

- **while** ((i = i + 1) < max)
 doSideEffect("foo")

4.13 Do-While Loop

A do-while loop **do** expression **while** (predicate) is used to execute a certain expression unless the predicate is evaluated to **false**. The difference to the while loop (§4.12) is that the execution starts by executing the block once before evaluating the predicate for the first time. The return type of a do-while loop is **void**.

4.13.1 Examples

- ```
do {
 doSideEffect("foo");
} while (true)
```
- ```
do doSideEffect("foo") while ((i=i+1)<max)
```

4.14 Return Expression

Although an explicit return is often not necessary, it is supported. In a closure for instance a return expression is always implied if the expression itself is not of type **void**. Anyway you can make it explicit:

```
listOfStrings.map(e| {  
    if (e==null)  
        return "NULL"  
    e.toUpperCase  
})
```

4.15 Throwing Exceptions

Like in Java it is possible to throw `java.lang.Throwable`. The syntax is exactly the same as in Java.

```
{  
    ...  
    if (myList.isEmpty)  
        throw new IllegalArgumentException("the list must not be empty")  
    ...  
}
```

4.16 Try, Catch, Finally

The try-catch-finally expression is used to handle exceptional situations. You are not forced to declare checked exceptions, if you do not catch checked exceptions they are rethrown in a wrapping runtime exception. Other than that the syntax again is like the one known from Java.

```
try {  
    throw new RuntimeException()  
} catch (NullPointerException e) {  
    // handle e  
} finally {  
    // do stuff  
}
```

4.17 Template Expressions

Templates allow for readable string concatenation, which is the main thing you do when writing a code generator. Let us have a look at an example of how a typical method with template expressions looks like:

```
toClass(Entity e) '''
    package «e.packageName»;

    «placImports»

    public class «e.name» «IF e.extends!=null»extends «e.extends»«ENDIF» {
        «FOR e.members»
            «member.toMember»
        «ENDFOR»
    }
'''
```

If you are familiar with Xpand, you will notice that it is exactly the same syntax. The difference is, that the template syntax is actually an expression, which means it can occur everywhere where an expression is expected. For instance in conjunction the powerful switch expression (§4.9):

```
toMember(Member m) {
    switch m {
        Field : '''private «m.type» «m.name» ;'''
        Method case isAbstract : ''' abstract «...'''
        Method : ''' ..... '''
    }
}
```

4.17.1 Conditions in Templates

There is a special **IF** to be used within templates which is identical in syntax and meaning to the old **IF** from Xpand. Note that you could also use the if expression, but since it has not an explicit terminal token, it is not as readable in that context.

4.17.2 Loops in Templates

Also the **FOR** statement is available and can only be used in the context of a template. It also supports the **SEPARATOR** from Xpand. In addition, a **BEFORE** expression can be defined that is only evaluated if the loop is at least evaluated once before the very first iteration. Consequently **AFTER** is evaluated after the last iteration if there is any element.

4.17.3 Typing

The rich string is translated to an efficient string concatenation and the return type of a rich string is CharSequence which allows room for efficient implementation.

4.17.4 White Space Handling

One of the key features of templates is the smart handling of white space in the template output. The white space is not written into the output data structure as is but preprocessed. This allows for readable templates as well as nicely formatted output. This can be achieved by applying three simple rules when the rich string is evaluated.

1. An evaluated rich string as part of another string will be prefixed with the current indentation of the caller before it is inserted into the result.
2. Indentation in the template that is relative to a control structure will not be propagated to the output string. A control structure is a **FOR**-loop or a condition (**IF**) as well as the opening and closing marks of the rich string itself. The indentation is considered to be relative to such a control structure if the previous line ends with a control structure followed by optional white space. The amount of white space is not taken into account but the delta to the other lines.
3. Lines that do not contain any static text which is not white space but do contain control structures or invocations of other templates which evaluate to an empty string, will not appear in the output.

The behavior is best described with a set of examples. The following table assumes a data structure of nested nodes.

```
class Template {  
    print(Node n) '''  
        node «n.name» {}  
    ...  
}
```

```
node NodeName{}
```

The indentation before `node «n.name»` will be skipped as it is relative to the opening mark of the rich string and thereby not considered to be relevant for the output but only for readability of the template itself.

```
class Template {  
    print(Node n) '''  
        node «n.name» {  
            «IF hasChildren»  
                «n.children*.print»  
            «ENDIF»  
        }  
    ...  
}
```

```
node Parent{  
    node FirstChild {  
    }  
    node SecondChild {  
        node Leaf {  
        }  
    }  
}
```

As in the previous example, there is no indentation on the root level for the same reason. The first nesting level has only one indentation level in the output. This is derived from the indentation of the **IF** `hasChildren` condition in the template which is

nested in the node. The additional nesting of the recursive invocation `children*.print` is not visible in the output as it is relative the the surrounding control structure. The line with **IF** and **ENDIF** contain only control structures thus they are skipped in the output. Note the additional indentation of the node *Leaf* which happens due to the first rule: Indentation is propagated to called templates.

List of External Links

http://java.sun.com/docs/books/jls/third_edition/html/conversions.html#5.5
<http://projectlombok.org/features/SneakyThrows.html>
http://java.sun.com/docs/books/jls/third_edition/html/conversions.html
http://java.sun.com/docs/books/jls/third_edition/html/j3TOC.html

Todo list

describe how to get it 4